
Learning-Rate Schedules Are Not Adiabatic: A Rate-Dependent Correction for Loss-Curve Prediction

Jiaju Wu

Peking University

<https://github.com/wjjpku/DL-final>

Abstract

State-of-the-art laws for predicting the pretraining loss curve under a learning-rate (LR) schedule—the Multi-Power Law (MPL) and the Functional Scaling Law (FSL)—are *adiabatic*: they predict the quasi-static loss the schedule would reach at equilibrium. MPL does use the LR decrements (a convolution of $\Delta\eta_k$ with a saturating kernel), but its only handle on the *rate* of decay is the single empirical exponent γ ; calibrated on slow schedules, this cannot capture the transient by which the loss *lags* its equilibrium when the LR is annealed fast. The result is a systematic *positive* residual on *fast* (warmup-stable-decay, WSD) decays, which we identify as a non-adiabatic relaxation lag. The classical SGD relaxation time governing this lag, $\tau \sim 1/(\eta\lambda)$, is well known; our contribution is to (i) show it governs the residual of adiabatic loss-curve laws, (ii) derive the resulting correction in closed form from AdamW’s preconditioned second-moment dynamics—a convolution of the per-step LR decrements with a cumulative-LR (S-time) exponential kernel, DropRelaxS, whose amplitude equals the exactly-identified equilibrium-floor sensitivity $dL_{\text{eq}}/d\eta$ —and (iii) turn it into a usable predictor. We confirm in LLM pretraining that the residual relaxes with the predicted $\tau \propto 1/\eta$ (measured pooled exponent $p = 1.00 \pm 0.18$), which also self-consistently establishes the noise-dominated regime the derivation assumes. Because the kernel’s rate constant is an *approximately scale-invariant* preconditioned curvature, the correction can be *predicted, not fitted*, and transferred across model size (same architecture): with zero fitting on the target curve it improves held-out sharp-decay loss prediction by 28–56% (44% on average). The correction is *complementary* to MPL’s empirical exponent γ (which tunes the adiabatic annealing kernel): it adds the non-adiabatic transient and improves a full- γ MPL. An independent reproduction on a separately-trained $\sim 10\text{M}$ transformer and a from-scratch AdamW simulation confirms the effect (the fast sweep ends higher *despite* 37% more cumulative LR) and the amplitude identity (within 3%). We are candid about scope: the gain is in the regime where WSD curves are unavailable; given them, a refit MPL absorbs the residual.

1 Introduction

The learning-rate schedule shapes the *entire* loss trajectory of LLM pretraining [11, 12, 13]. *Schedule-aware loss-curve prediction* asks: given $\{\eta_t\}$, predict $L(t)$, so that schedules can be compared and optimized *before* spending compute. The defining test is cross-schedule generalization—fit on cosine schedules, predict the unseen WSD family. The state of the art is the **Multi-Power Law** (MPL) [2],

with a matching kernel-regression theory, the **Functional Scaling Law** (FSL) [3]:

$$L(t) = L_0 + A S(t)^{-\alpha} + B \sum_{k \leq t} \Delta \eta_k G(\eta_k^{-\gamma} (S(t) - S(k))), \quad G(x) = 1 - (1 + Cx)^{-\beta}, \quad (1)$$

where $S(t) = \sum_{i \leq t} \eta_i$ and $\Delta \eta_k = \eta_{k-1} - \eta_k$.

The gap: existing laws are adiabatic. These laws *do* use the LR decrements: MPL convolves $\Delta \eta_k$ with a saturating kernel G , and its factor $\eta_k^{-\gamma}$ is a rate-of-decay knob. But this single empirical exponent is their *only* handle on how *fast* the LR changes. Structurally they predict the *quasi-static* loss—the value the schedule would reach if the loss stayed at equilibrium; the river-valley analysis [4] makes this explicit, modelling the noise penalty as set *instantaneously* by the current LR. We call such laws *adiabatic*. Calibrated on slow schedules, γ does not extrapolate to fast WSD decays, where the loss genuinely lags equilibrium—leaving a structured residual (Fig. 1) that no single rate exponent captures across schedules. Our correction is a derived transient *complementary* to γ , which tunes the adiabatic kernel.

Contributions.

- We show that adiabatic loss-curve laws leave a systematic *positive* residual on fast decays, and identify it as a non-adiabatic relaxation lag (§2).
- From AdamW’s preconditioned second-moment dynamics we *derive* the leading non-adiabatic correction in closed form, the rate-dependent term DropRelaxS, with amplitude equal to the exactly-identified $dL_{\text{eq}}/d\eta$ (§3). The relaxation time it uses is the classical $\tau \sim 1/(\eta\lambda)$; our point is that it governs the *loss-curve-law residual*, not that the timescale is new.
- We confirm in LLM pretraining that the residual relaxes with $\tau \propto 1/\eta$ ($p = 1.00 \pm 0.18$), self-consistently establishing its noise-dominated regime (§4).
- The kernel’s rate constant is an approximately scale-invariant preconditioned curvature, enabling a *parameter-free, cross-scale* prediction that improves held-out sharp-decay loss by 44% with no target-curve fitting (§4).
- We give a fair-baseline diagnostic separating a genuine formula gain from an information gap, clarifying exactly when the correction helps (§4,6).

All experiments use the public MPL loss curves (Llama-2-style 25/100/400M, 9 schedules each) [2]; a faithfully reproduced MPL (test MAE 0.0054/0.0050/0.0085, matching [2]) is the baseline throughout.

2 The non-adiabatic residual

We take a *well-fit* MPL (the published parameters [2]) and examine its prediction residual $r(t) = L_{\text{true}}(t) - L_{\text{MPL}}(t)$ on held-out schedules. (Isolating the non-adiabatic term requires a well-fit MPL: a baseline that already captures the *adiabatic* part cleanly, so the remainder is the lag and not a backbone mis-fit; a cosine-only fit conflates the two and gives a far noisier residual.) The residual has a sharp signature (Fig. 1): it is ≈ 0 throughout the long stable phase and on gradual cosine decay, but jumps *positive* exactly when the LR decays steeply, and is largest—and *undamped at the curve’s end*—for the sharpest decays (`wsd`, `wsdld`). The end-of-curve lag is $+(6\text{--}15) \times 10^{-3}$ on WSD/WSDLD versus $\lesssim 4 \times 10^{-3}$ on cosine. Crucially, a gradual cosine and a sharp WSD decay reach the *same* final LR, yet only the fast one lags: same destination, faster sweep, bigger lag—the textbook signature of a rate-dependent (non-adiabatic) effect, and exactly the dependence an adiabatic law cannot represent.

3 Theory: the leading non-adiabatic correction

MPL is the adiabatic limit. At a fixed LR the loss relaxes toward a quasi-static value $L_{\text{eq}}(\eta)$ (backbone progress plus the equilibrium variance of the sharp “mountain” directions of the river-valley landscape [4]). On schedules that vary the LR slowly the loss tracks L_{eq} , so a **well-fit MPL approximates the adiabatic** L_{eq} (its kernel, in particular γ , is calibrated to reproduce the quasi-static component). The residual on *fast* held-out decays is the non-adiabatic remainder.




Figure 1: The well-fit MPL prediction residual on the held-out *wsd* curve (solid) is ≈ 0 during the stable phase and jumps up at the step-20000 decay. The derived rate-dependent term κ DropRelaxS (dashed) captures it ($R^2=0.40/0.80/0.87$ at 25/100/400M; the signal, and the fit, grow with model size).

Derivation (AdamW). In the Hessian eigenbasis, mode i has curvature λ_i and gradient-noise variance s_i^2 . With $\beta_2 \rightarrow 1$ the Adam preconditioner sits at $v_i^* = \mathbb{E}[g_i^2] = \lambda_i^2 V_i + s_i^2$, and in the noise-dominated mountain directions $v_i^* \approx s_i^2$, so Adam reduces to SGD with a *preconditioned* step $\eta_i^{\text{eff}} = \eta/s_i$ [7, 9]. The second moment $V_i = \mathbb{E}[\delta_i^2]$ then obeys the *exact* recursion $V_i(t+1) = (1 - \eta_i^{\text{eff}} \lambda_i)^2 V_i(t) + (\eta_i^{\text{eff}})^2 s_i^2$, with fixed point $V_i^*(\eta)$ and the classical relaxation rate $2\eta_i^{\text{eff}} \lambda_i$, i.e. $\tau_i \sim 1/(\eta \lambda_i/s_i)$ [6]. Writing the deviation $\Delta_i = V_i - V_i^*(\eta_t)$ and using $(1 - \eta_i^{\text{eff}} \lambda_i)^2 \approx e^{-2\eta_i^{\text{eff}} \lambda_i}$ (valid because the lag lives on *slow* modes, $\eta \lambda_i/s_i \ll 1$, verified post hoc), the linear recursion solves to a cumulative-LR (S-time) exponential:

$$\Delta_i(t) = \frac{dV_i^*}{d\eta} \sum_{t'} e^{-2(\lambda_i/s_i)(S_t - S_{t'})} \text{drop}_{t'}, \quad \text{drop}_{t'} = (\eta_{t'-1} - \eta_{t'})_+. \quad (2)$$

The loss lag $\Delta L = \frac{1}{2} \sum_i \lambda_i \Delta_i$ is therefore a w_i -weighted *mixture* of S-time exponentials with rates $2\lambda_i/s_i$, and the total amplitude is an *exact identity*:

$$\Delta L(t) = \sum_i w_i [e^{-2(\lambda_i/s_i)(S - S')} \otimes \text{drop}](t), \quad \sum_i w_i = \frac{dL_{\text{eq}}}{d\eta}. \quad (3)$$

Collapsing the spectrum to one effective rate gives the operational law (one shape constant λ_{slow} , one calibration $c \leq 1$):

$$L(t) = L_{\text{MPL}}(t) + c \cdot \eta_{\text{peak}} \frac{dL_{\text{eq}}}{d\eta} \cdot \underbrace{\sum_{t'} e^{-\lambda_{\text{slow}}(S_t - S_{t'})} \text{drop}_{t'}}_{\text{DropRelaxS}_{\lambda_{\text{slow}}}(t)} \quad (4)$$

What the derivation predicts (all tested in §4). (i) *Sign*: $dV_i^*/d\eta > 0$, so $\Delta L > 0$ (loss above MPL). (ii) *Rate*: the per-step rate is $2(\eta/s_i)\lambda_i \propto \eta$, hence $\tau \propto 1/\eta$; a self-consistent treatment of the preconditioner gives $\tau \propto \eta^{-p}$ with $p=1$ in the noise-dominated regime, which the data occupy. (A putative signal-dominated $p=2/3$ does not arise in practice—forcing signal dominance still yields $p \approx 1$, §5—so we do not rely on it.) (iii) *Amplitude*: $\sum w_i = dL_{\text{eq}}/d\eta$, measurable from the noise floor. (iv) *Shape*: $\lambda_{\text{slow}} = 2(\lambda/s)_{\text{eff}}$ is a *preconditioned* curvature—the quantity Adam stabilizes [9]—so it is roughly scale/architecture invariant. (v) *Spectral mixture*: Eq. (3) is a sum of exponentials, so a two-exponential kernel should beat one.

4 Experiments

We freeze the published MPL and compare it to MPL+Eq. (4); the only addition is the rate-dependent term. Unless noted $\lambda_{\text{slow}}=10$.




Figure 2: Loss-relaxation time τ measured from `wsdcon` decay transients vs. stage-2 LR. The 100/400M curves track the predicted slope -1 ($\tau \propto 1/\eta$); pooled $p = 1.00 \pm 0.18$.

(i) The rate prediction $\tau \propto 1/\eta$ holds (Fig. 2). The two-stage `wsdcon` curves step the LR down at step 8000 to a constant $\eta \in \{3, 9, 18\} \times 10^{-5}$, giving a clean relaxation window. Fitting the residual transient $r(t) \propto e^{-(t-8000)/\tau}$ yields τ that scales as η^{-p} with *pooled* $p = 1.00 \pm 0.18$ (per scale $-1.06/-0.94/-0.51$ at 100/400/25M). This is the classical $\tau \sim 1/(\eta\lambda)$ confirmed for the loss-curve-law residual in LLM pretraining, and—since $p=1$ is the noise-dominated value—it self-consistently justifies the $v_i^* \approx s_i^2$ step of the derivation. The implied rate constant $\lambda_{\text{slow}} = 1/(\tau\eta) \approx 10$ is scale-invariant; its modes have $\eta_{\text{peak}}\lambda_{\text{eff}} \approx 2 \times 10^{-3} \ll 2$ (deep in the slow regime, far from the edge of stability—unlike γ).

(i') The same scaling holds in a from-scratch AdamW simulation. To check that $\tau \propto 1/\eta$ is the mechanism and not a data artefact, we simulate AdamW on a noisy quadratic from scratch, equilibrate at a high LR, step it down, and fit the loss-relaxation time (`repro/beta2_ablation.py`). Sweeping the post-step LR reproduces $\tau \propto 1/\eta$ with log-log slope $p = 1.01$ —the mechanism holds in a controlled optimiser, independent of the public curves.

(ii) The residual is the derived term. Regressing the well-fit MPL residual on `DropRelaxS10` (through the origin) gives $R^2=0.40/0.80/0.87$ at 25/100/400M (Fig. 1); a 2-exponential kernel improves R^2 by 0.06–0.07, confirming the predicted spectral mixture (Eq. 3). The form is not arbitrary: in a fair few-shot comparison the derived S-time exponential of decrements beats a step-time kernel (-21% vs. -11% held-out MAE) and a memoryless floor (-3%), confirming that the cumulative-LR “time” and the decaying memory both matter (`repro/generic_kernel_compare.py`). (On a *single* decay curve the finite- λ_{slow} kernel is numerically close to the $\lambda_{\text{slow}} \rightarrow 0$ adiabatic limit, so the rate constant is pinned only *across* schedules; §5.)

(iii) The amplitude is predictable from the noise floor. Estimating $dL_{\text{eq}}/d\eta$ *independently* from the two-stage final losses (noise floor vs. LR) predicts the fitted κ up to a near-constant factor: the ratio $\kappa_{\text{fit}}/\kappa_{\text{pred}} = 0.43/0.51/0.60$ has coefficient of variation 14%, i.e. $c \approx 0.5$ is universal (the < 1 is the spectral spread of Eq. 3).

(iv) Parameter-free, cross-scale prediction (Table 1). Because λ_{slow} and c are approximately scale-invariant across these (same-architecture) scales and $dL_{\text{eq}}/d\eta$ is read off the (cheap) noise floor, the whole correction can be *predicted with zero fitting on the target curve*: take $\lambda_{\text{slow}}=10$, c from a leave-one-scale-out average of the *other* scales, and $dL_{\text{eq}}/d\eta$ from the target’s own two-stage probes (distinct curves from the test). Applied to the held-out `wsd/wsd1d` decays it cuts MAE by 28–56% (44% average, wins 6/6)—small-model constants plus a cheap probe predict the large-model sharp-decay curve $\sim 2\times$ better than MPL transfer.

(v) Honest scope: when it does *not* help. If instead one refits MPL on a few *full* WSD curves, MPL itself absorbs the residual (held-out MAE $0.0037 \rightarrow 0.0032$) and the explicit term adds nothing. So the residual is, for a data-rich MPL, an *information* gap, not a formula gap: the value of the correction

Table 1: Parameter-free cross-scale prediction: well-fit MPL vs. MPL+Eq. (4) with *no* fitting on the target curve. MAE on held-out sharp decays.

Scale	wsd		wsd+d		avg. Δ
	MPL	+ours	MPL	+ours	
25M	0.00341	0.00246	0.00314	0.00219	−29%
100M	0.00375	0.00164	0.00325	0.00161	−53%
400M	0.00470	0.00236	0.00385	0.00212	−47%
overall					−44% (6/6)




Figure 3: Residual vs. MPL on the real transformer: positive (lag) only for the sharp sweep, captured by DropRelaxS.

Table 2: Real ~ 10 M transformer: three schedules to the *same* final LR. The fast sweep ends highest despite the most cumulative LR.

schedule	final loss	cumulative LR S
cosine (gradual)	1.0490	8.20
wsd (gradual)	1.0482	9.22
wsd (sharp)	1.0593	12.59

is precisely the data-*poor* regime—cosine plus cheap probes, but no full WSD curves—which is the actual “predict before training” use case.

(vi) **Relation to γ : complementary, not a substitute.** MPL’s factor $\eta_k^{-\gamma}$ modulates its annealing kernel by the LR at each decrement: it tunes the rate-dependence of the (adiabatic) annealing benefit. Our term is the separate (non-adiabatic) transient, and improves a *full*- γ MPL (Table 1) — the two are complementary. Characterising their precise relationship is left to future work (§7).

5 Independent reproduction

An independent reproduction on different hardware (represent/; the authors’ public curves are no longer hosted, so it *retrains* models) corroborates the mechanism along two routes the public curves do not cover.

A real, independently-trained transformer (no fitting). A ~ 10 M-parameter Llama-2-style byte-level transformer was trained on WikiText with AdamW under three schedules reaching the *same* final LR ($2.5 \times 10^{-3} \rightarrow 2.5 \times 10^{-4}$) at different speeds (Table 2). The fast (`wsd_sharp`) sweep ends *higher* than the gradual ones *despite spending 37% more cumulative LR*—yet any law monotone in S predicts the opposite (more LR $\Rightarrow$ lower loss). With no fitting, this is the rate-dependent non-adiabatic lag, on a model and data we never used; the residual against a well-fit MPL is positive only for the sharp sweep and is captured by DropRelaxS (Fig. 3). Across schedules a finite- λ_{slow} kernel beats the adiabatic ($\lambda_{\text{slow}} \rightarrow 0$) baseline by $\Delta R^2 \approx +0.11$ held out.

From-scratch AdamW confirms the mechanism quantitatively. Simulating AdamW on a noisy quadratic (Monte-Carlo ground truth; closed-form theory used only as an independent baseline), the reproduction recovers $\tau \propto 1/\eta$ with $p = 0.971$ (log-log $R^2=0.9996$, robust to leave-one-out and seeds), τ independent of β_2 (CV 1.1%), the **amplitude identity** $dL_{\text{eq}}/d\eta = 3.075$ versus the closed form $\frac{1}{4} \sum_i s_i = 3.0$ (within 3%), and a two-exponential kernel that recovers the injected preconditioned rates $2\lambda/s$. The effect is falsifiable, not circular: a signal-dominated spectrum breaks $p=1$.

What does not transfer (consistent with our scope). At $\sim 10\text{M}$ the $\tau \propto 1/\eta$ relation itself is *not* resolved (the measured post-step relaxation is roughly flat over a $16\times$ LR range, audited as real and not a window artefact), and on a *single* decay curve the kernel is numerically degenerate with the adiabatic baseline, so the rate-dependence is identifiable only *across* schedules. This matches our stated scope: the effect is clean and present, while the quantitative law is a leading-order, data-scale-limited approximation whose constants (λ_{slow} , c , and the η_{peak} normalisation) remain measured rather than derived.

6 Related work

Loss-curve laws. Tissue et al. [1], MPL [2], and FSL [3] predict the trajectory from integrated-LR statistics (cumulative LR / intrinsic time / convolutional functionals). FSL and its follow-ups also derive optimal schedules [3]. All are adiabatic in our sense, and our rate-dependent term is the leading correction they omit; it is complementary to MPL’s exponent γ (which tunes the adiabatic kernel), not a replacement for it.

Cooldown dynamics and schedule geometry. The river-valley picture [4] explains qualitatively why decay “reveals progress”, and explicitly models the noise penalty as *instantaneous* in the current LR—the adiabatic assumption we refine. Dremov et al. [5] study the cooldown empirically (including the role of AdamW’s β_2 and the decay shape), framing it through a bias–variance trade-off rather than a rate-dependent lag. Late-decay analyses [11] and schedule-shape searches [10] study related questions (why a long high-LR phase helps; what optimal shapes look like) but do not model the loss’s rate dependence.

SGD relaxation physics. That a quadratic mode relaxes with $\tau \sim 1/(\eta\lambda)$ and that the SGD noise floor scales with η are classical, from noisy-quadratic and Langevin analyses of SGD [6]; adaptive methods inherit a preconditioned version [7], operating at an adaptive edge of stability [8, 9]. We do not claim these ingredients as new. Our contribution is the synthesis: identifying that adiabatic loss-curve laws miss exactly this relaxation lag, deriving the correction with an exactly-identified amplitude, and building a parameter-free cross-scale predictor from it.

7 Limitations

The gain is regime-specific (data-poor sharp decays); λ_{slow} and c are *measured*, not computed from first principles (this needs the noise spectrum $\{s_i\}$, which we do not have), and λ_{slow} is only *approximately* scale-invariant—it tracks the slowest preconditioned mode and drifts when the spectral *shape* changes (§5), so the cross-scale transfer is a leading-order, same-architecture approximation. The derivation also assumes a locally quadratic landscape with a static spectrum and diagonal, uncorrelated noise—the model-level assumptions MPL itself rests on. Finally, the quantitative $\tau \propto 1/\eta$ relation, clean in the controlled simulation, is *not* resolved on the $\sim 10\text{M}$ transformer (§5); confirming it on larger real models, and characterising the precise relationship to γ , are the main open items.

8 Conclusion

Existing loss-curve laws are adiabatic approximations. We identified, derived from AdamW’s preconditioned dynamics, and validated the leading non-adiabatic correction: a rate-dependent term whose relaxation time follows the classical $\tau \propto 1/\eta$ ($p = 1.00 \pm 0.18$, confirmed for the loss-curve residual) and whose scale-invariant shape enables parameter-free cross-scale prediction (−44% on

held-out sharp decays). The loss curve depends not only on how much LR has been spent, but on how fast it is being spent.

References

- [1] H. Tissue, V. Wang, and L. Wang. Scaling law with learning rate annealing. *arXiv:2408.11029*, 2024.
- [2] K. Luo, H. Wen, S. Hu, et al. A multi-power law for loss curve prediction across learning rate schedules. *ICLR; arXiv:2503.12811*, 2025.
- [3] B. Li, F. Chen, Z. Huang, L. Wang, and L. Wu. Functional scaling laws in kernel regression: loss dynamics and learning rate schedules. *NeurIPS; arXiv:2509.19189*, 2025.
- [4] K. Wen, Z. Li, T. Ma, et al. Understanding warmup-stable-decay learning rates: a river valley loss landscape perspective. *arXiv:2410.05192*, 2024.
- [5] A. Dremov, A. Hägele, A. Kosson, and M. Jaggi. Training dynamics of the cooldown stage in warmup-stable-decay learning rate scheduler. *TMLR; arXiv:2508.01483*, 2025.
- [6] G. Zhang, L. Li, Z. Nado, et al. Which algorithmic choices matter at which batch sizes? Insights from a noisy quadratic model. *NeurIPS*, 2019.
- [7] S. Malladi, K. Lyu, A. Panigrahi, and S. Arora. On the SDEs and scaling rules for adaptive gradient algorithms. *NeurIPS; arXiv:2205.10287*, 2022.
- [8] J. Cohen, S. Kaur, Y. Li, J. Z. Kolter, and A. Talwalkar. Gradient descent on neural networks typically occurs at the edge of stability. *ICLR*, 2021.
- [9] J. Cohen, B. Ghorbani, S. Krishnan, et al. Adaptive gradient methods at the edge of stability. *arXiv:2207.14484*, 2022.
- [10] H. Naganuma et al. What do near-optimal learning rate schedules look like? *arXiv:2603.10301*, 2026.
- [11] J. Kaplan, S. McCandlish, T. Henighan, et al. Scaling laws for neural language models. *arXiv:2001.08361*, 2020.
- [12] J. Hoffmann, S. Borgeaud, A. Mensch, et al. Training compute-optimal large language models. *NeurIPS*, 2022.
- [13] S. Hu, Y. Tu, X. Han, et al. MiniCPM: unveiling the potential of small language models. *COLM*, 2024.